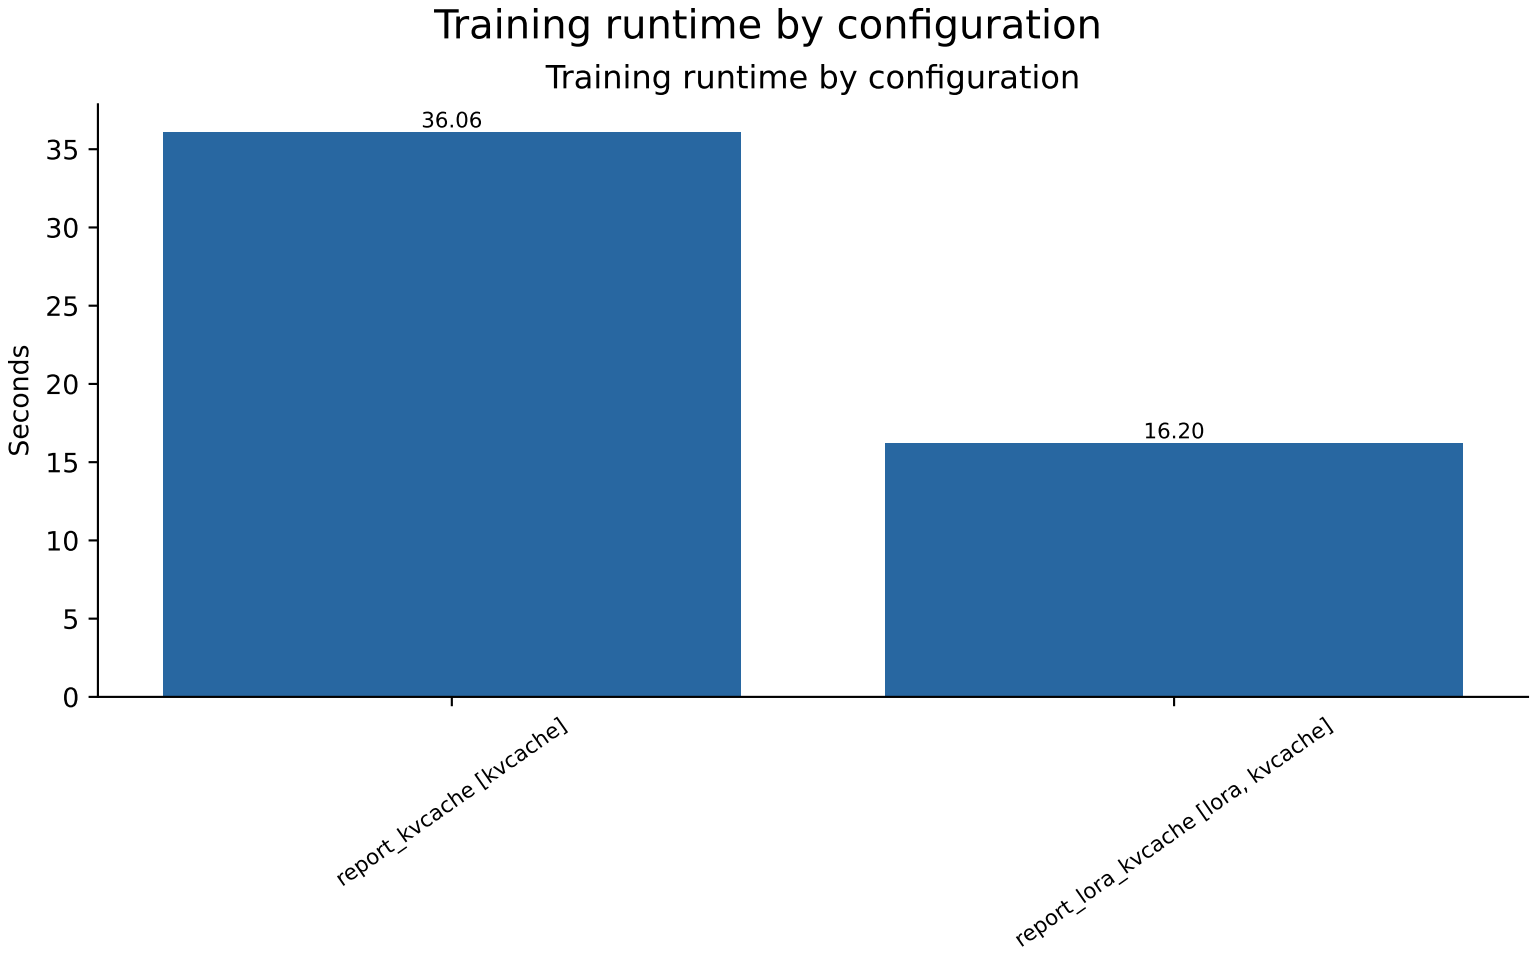
Training runtime by configuration
Training runtime by configuration
Seconds
0
5
10
15
20
25
30
35
36.06
16.20
report_kvcache [kvcache]
report_lora_kvcache [lora, kvcache]

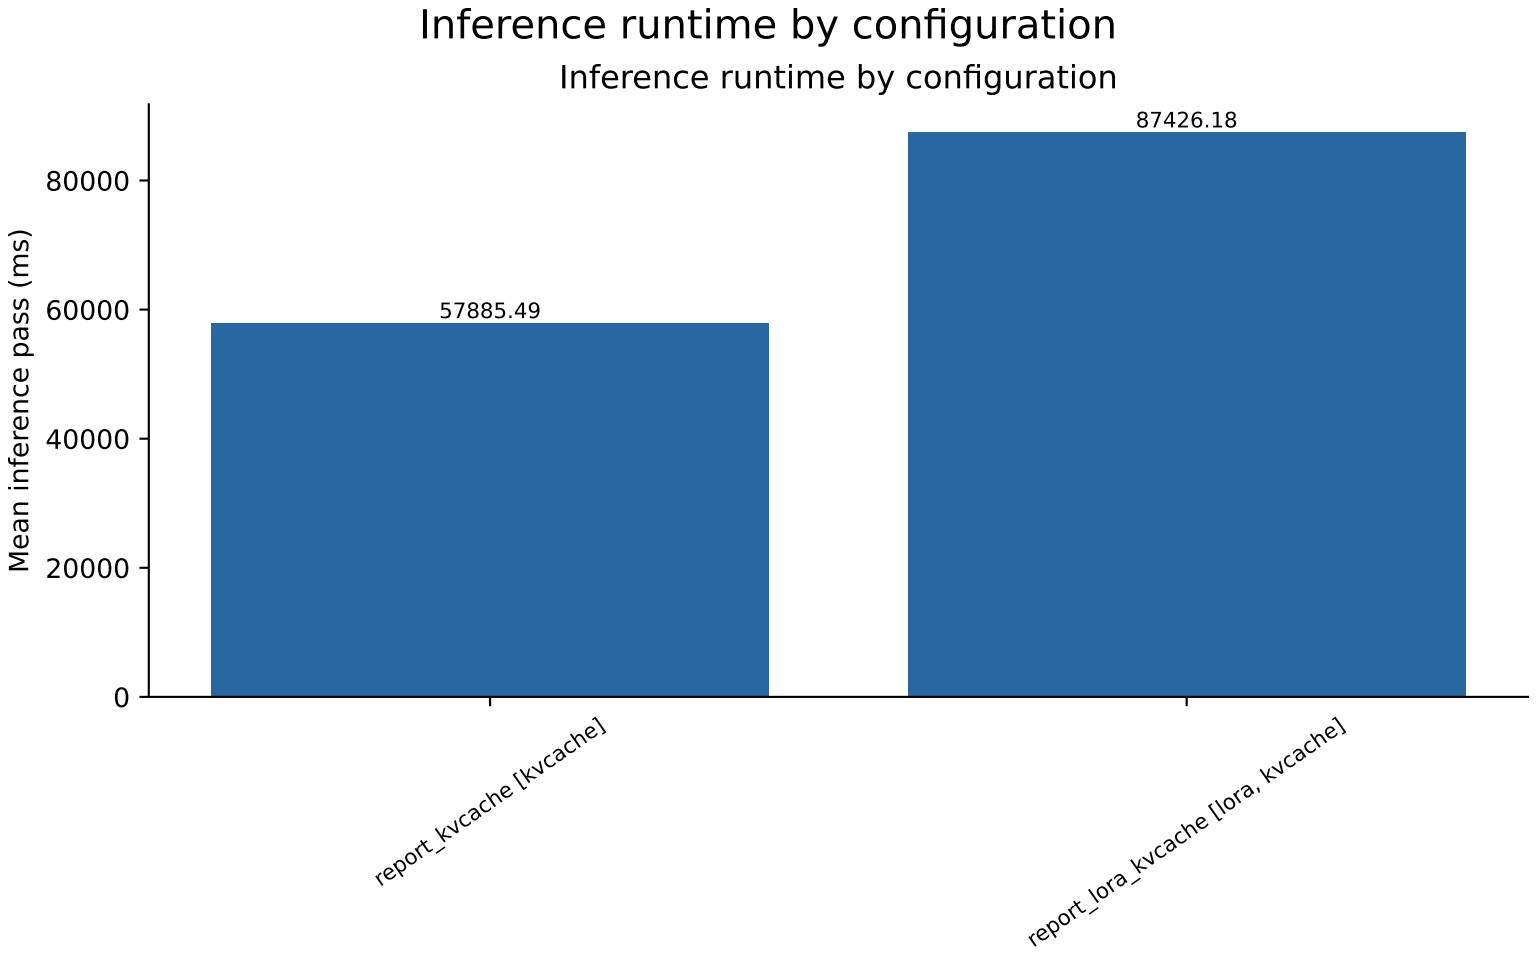
Inference runtime by configuration
Inference runtime by configuration
Mean inference pass (ms)
0
20000
40000
60000
80000
57885.49
87426.18
report_kvcache [kvcache]
report_lora_kvcache [lora, kvcache]

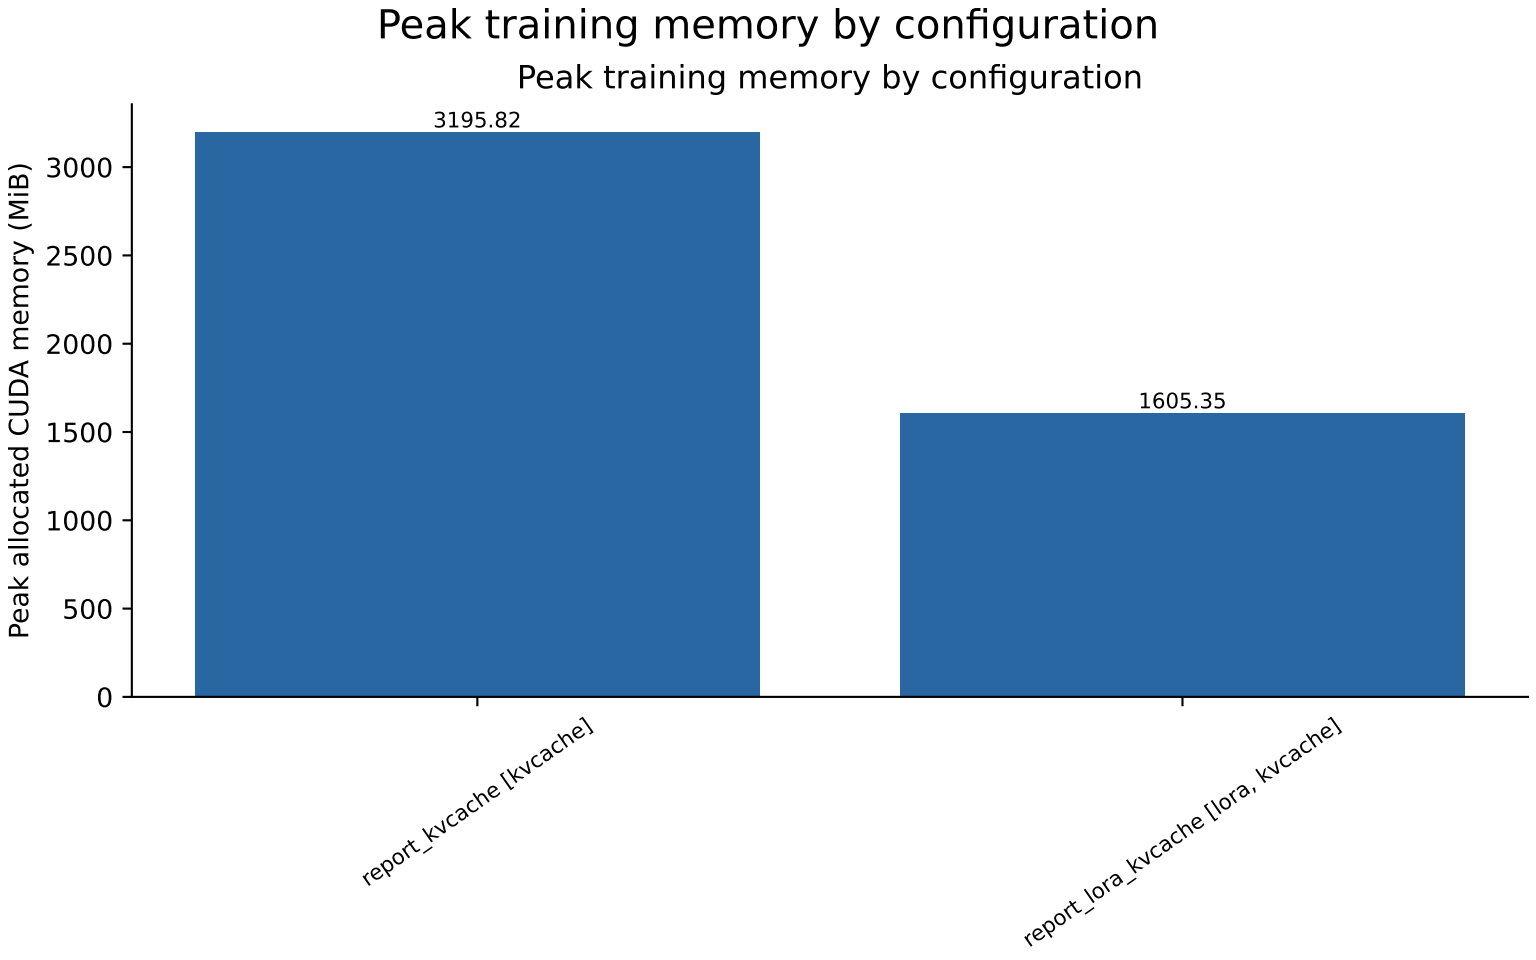
Peak training memory by configuration
Peak training memory by configuration
Peak allocated CUDA memory (MiB)
0
500
1000
1500
2000
2500
3000
3195.82
1605.35
report_kvcache [kvcache]
report_lora_kvcache [lora, kvcache]

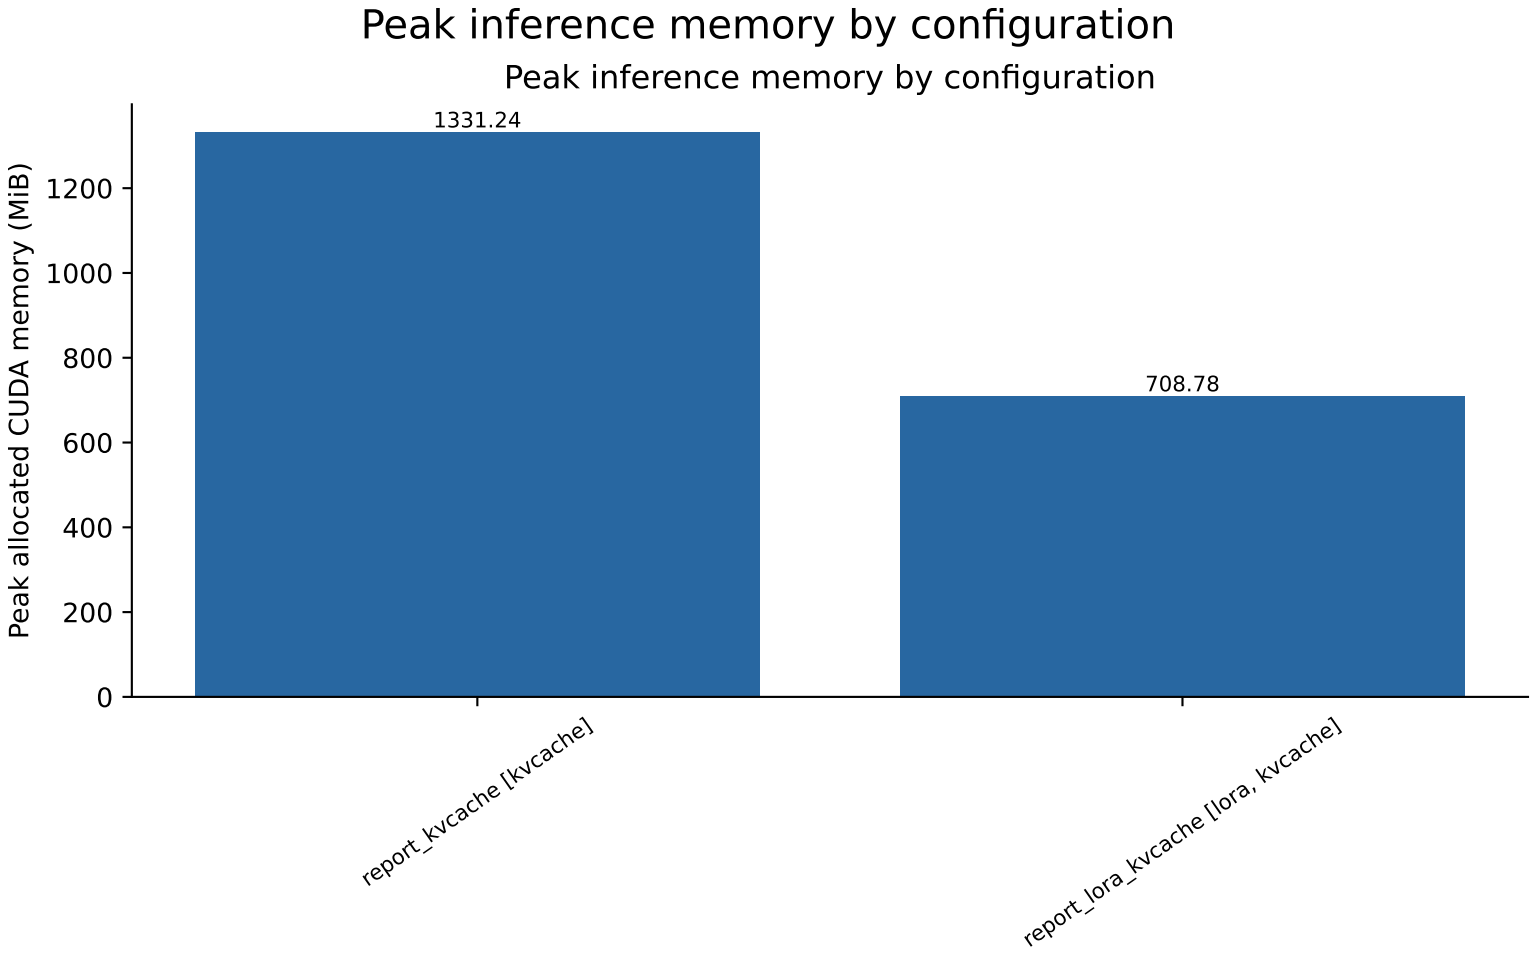
Peak inference memory by configuration
Peak inference memory by configuration
Peak allocated CUDA memory (MiB)
0
200
400
600
800
1000
1200
1331.24
708.78
report_kvcache [kvcache]
report_lora_kvcache [lora, kvcache]

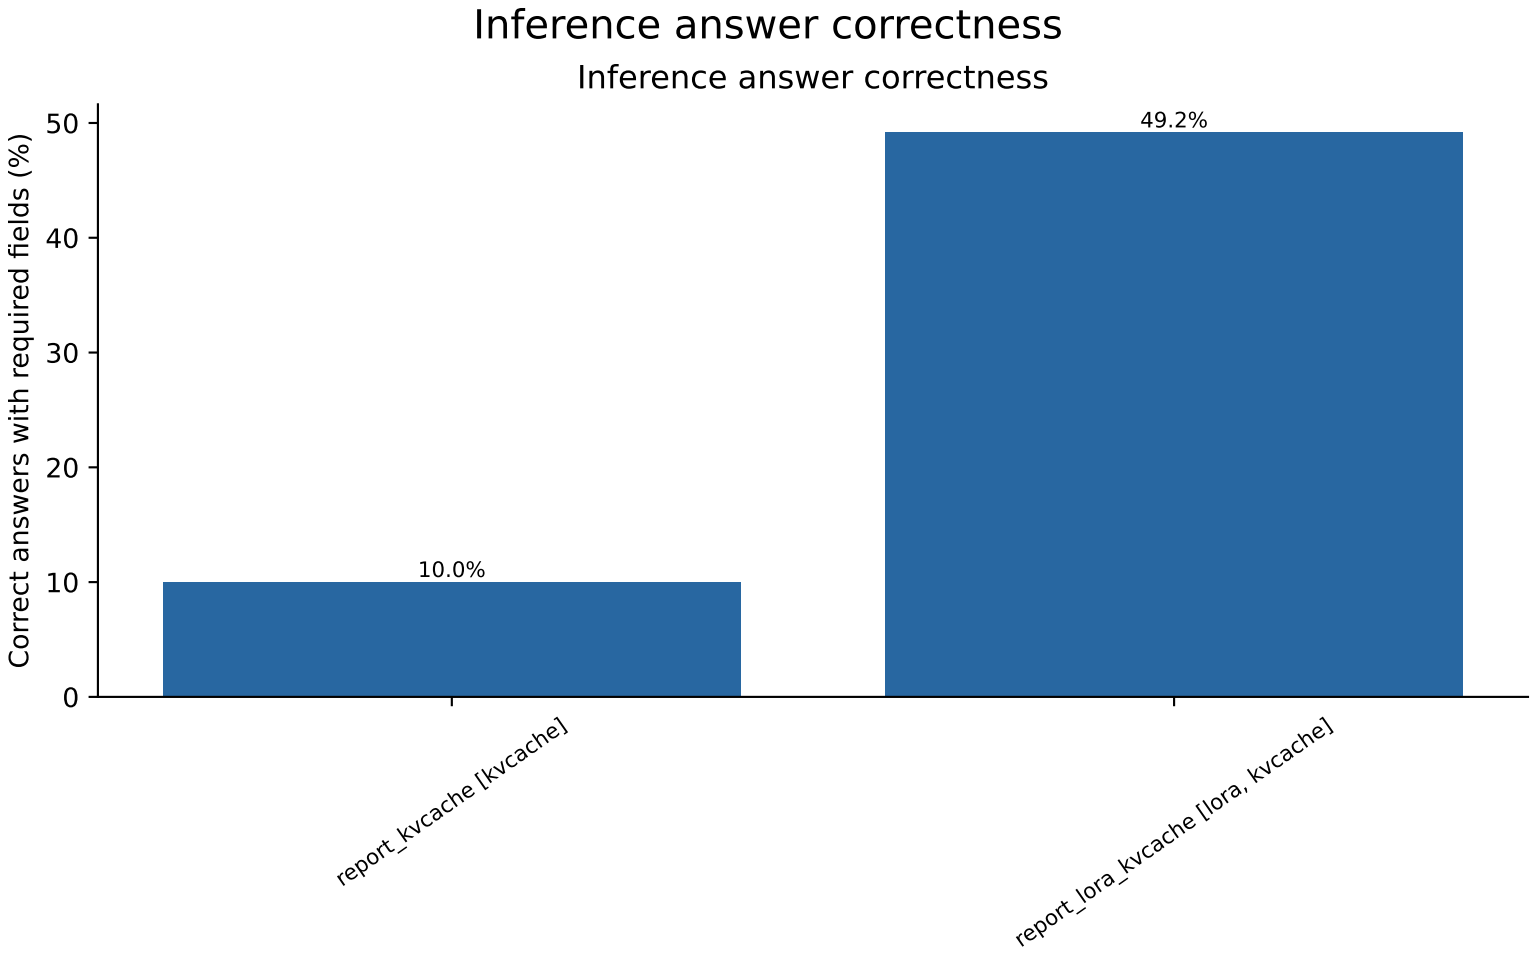
Inference answer correctness
Inference answer correctness
Correct answers with required fields (%)
0
10
20
30
40
50
10.0%
49.2%
report_kvcache [kvcache]
report_lora_kvcache [lora, kvcache]

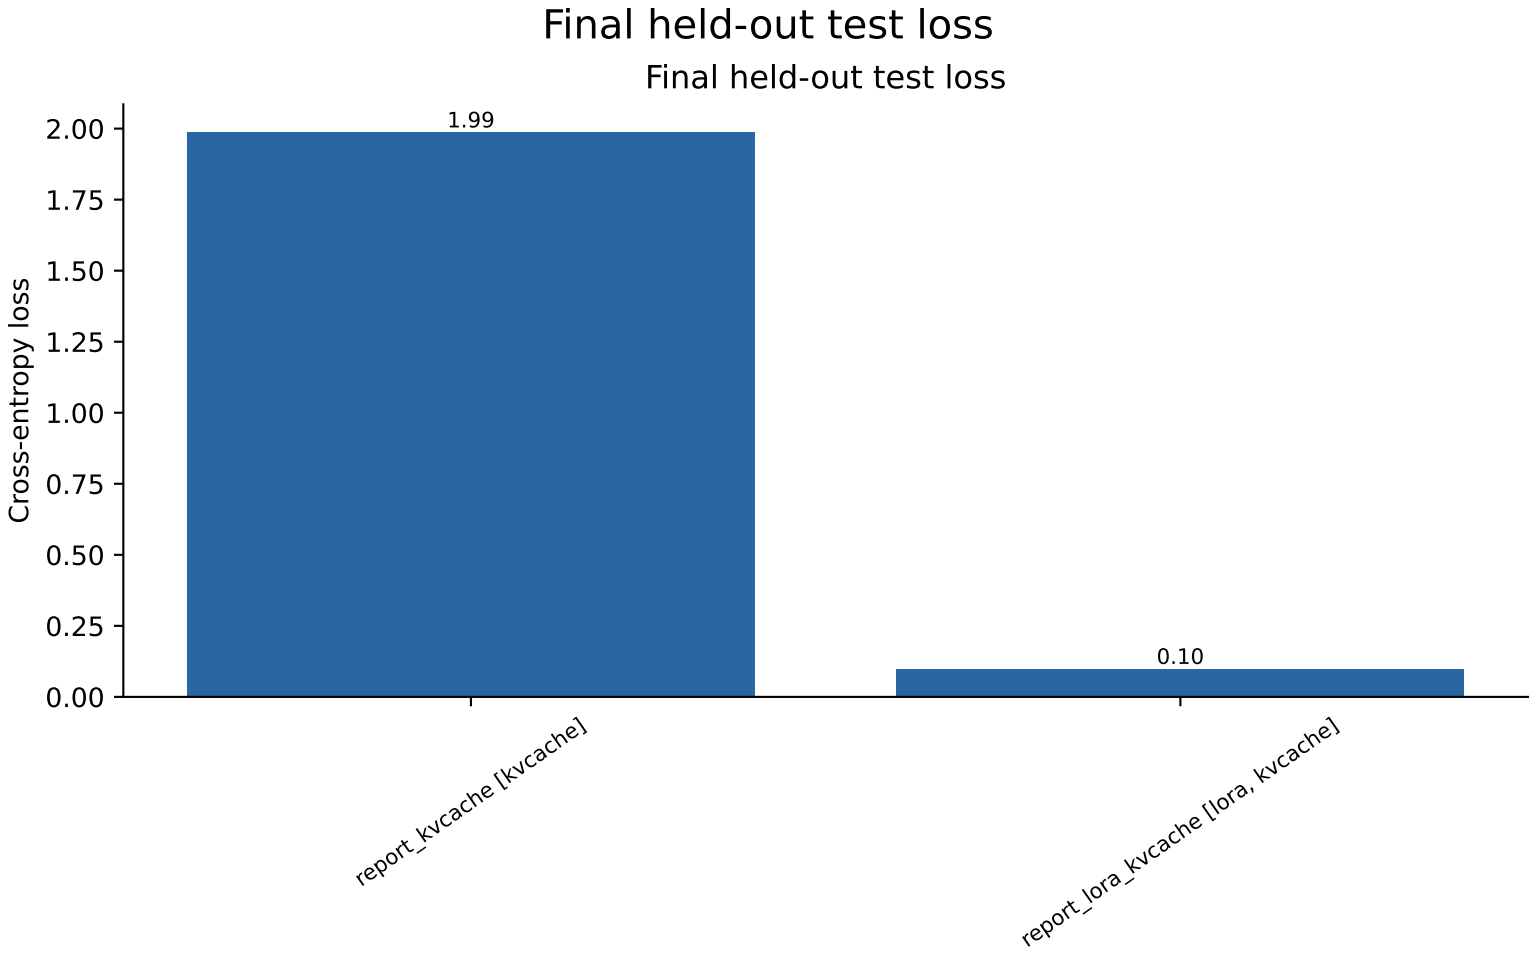
Final held-out test loss
Final held-out test loss
Cross-entropy loss
2.00
1.75
1.50
1.25
1.00
0.75
0.50
0.25
0.00
1.99
0.10
report_kvcache [kvcache]
report_lora_kvcache [lora, kvcache]
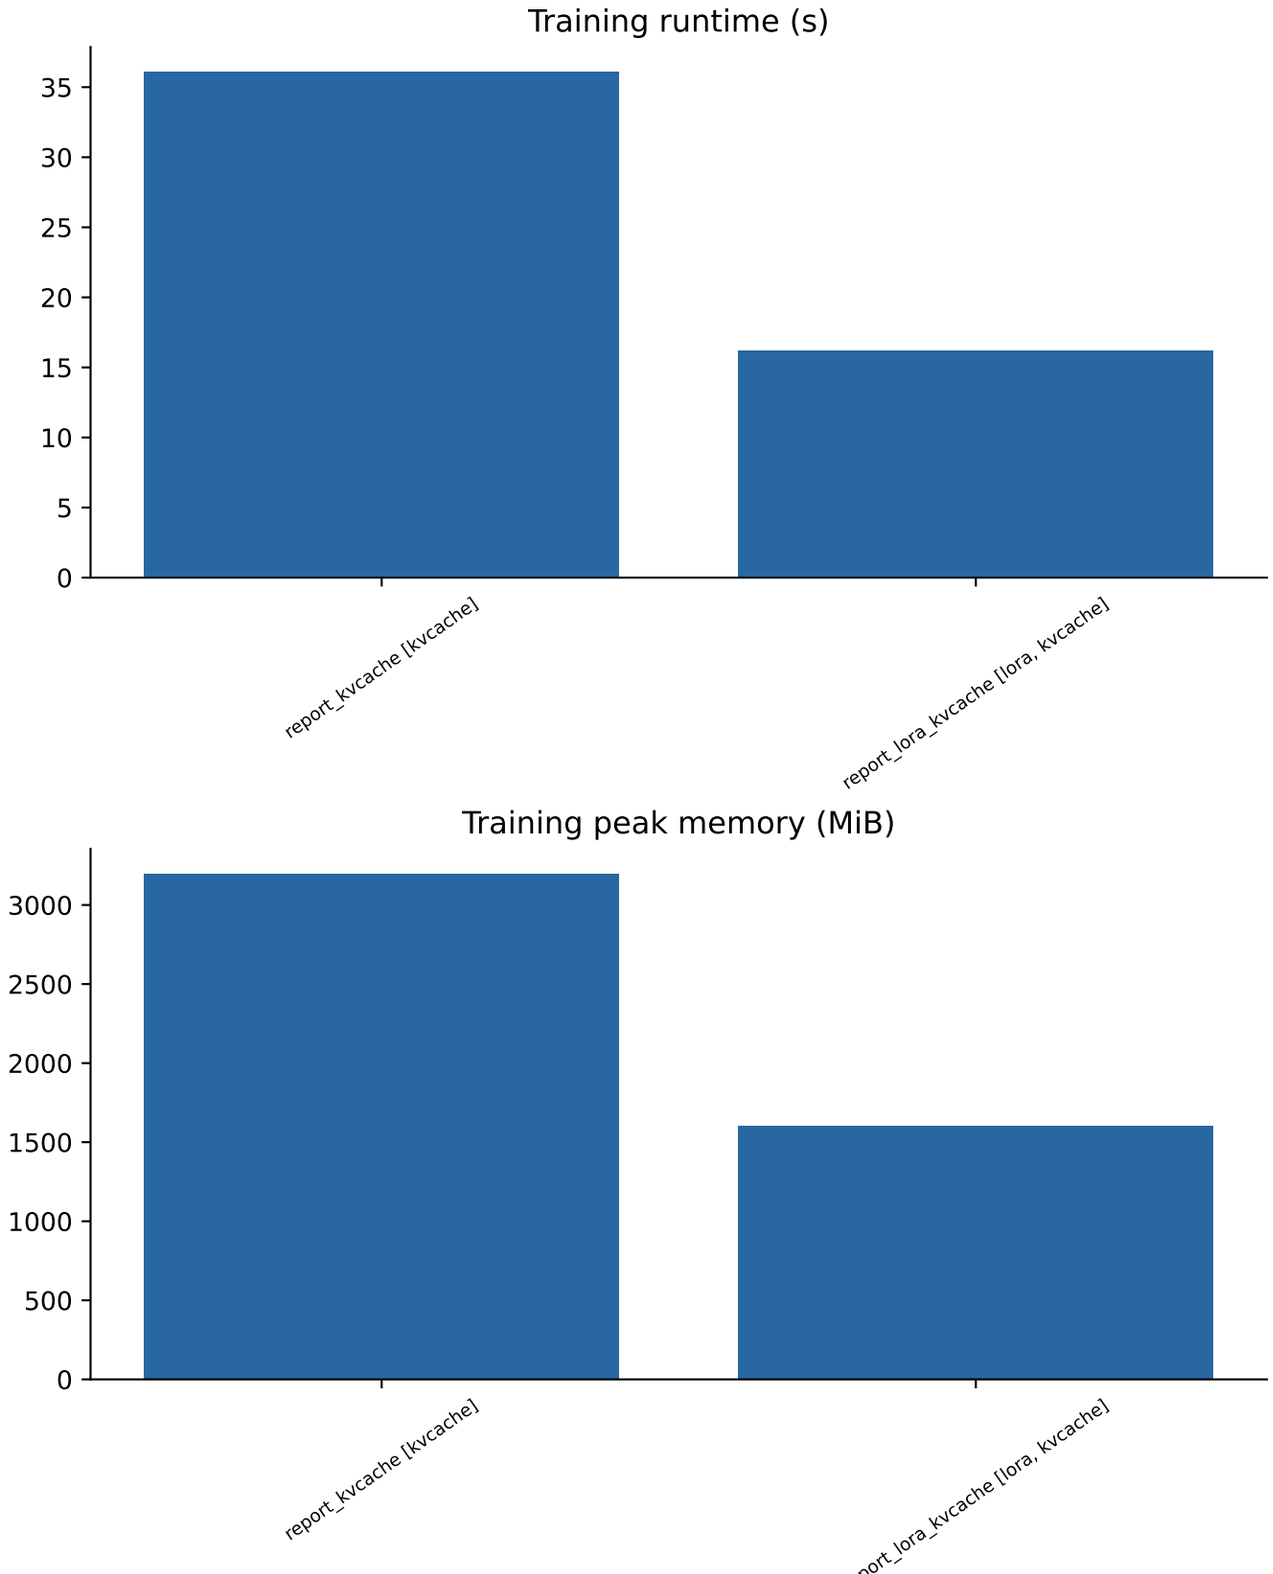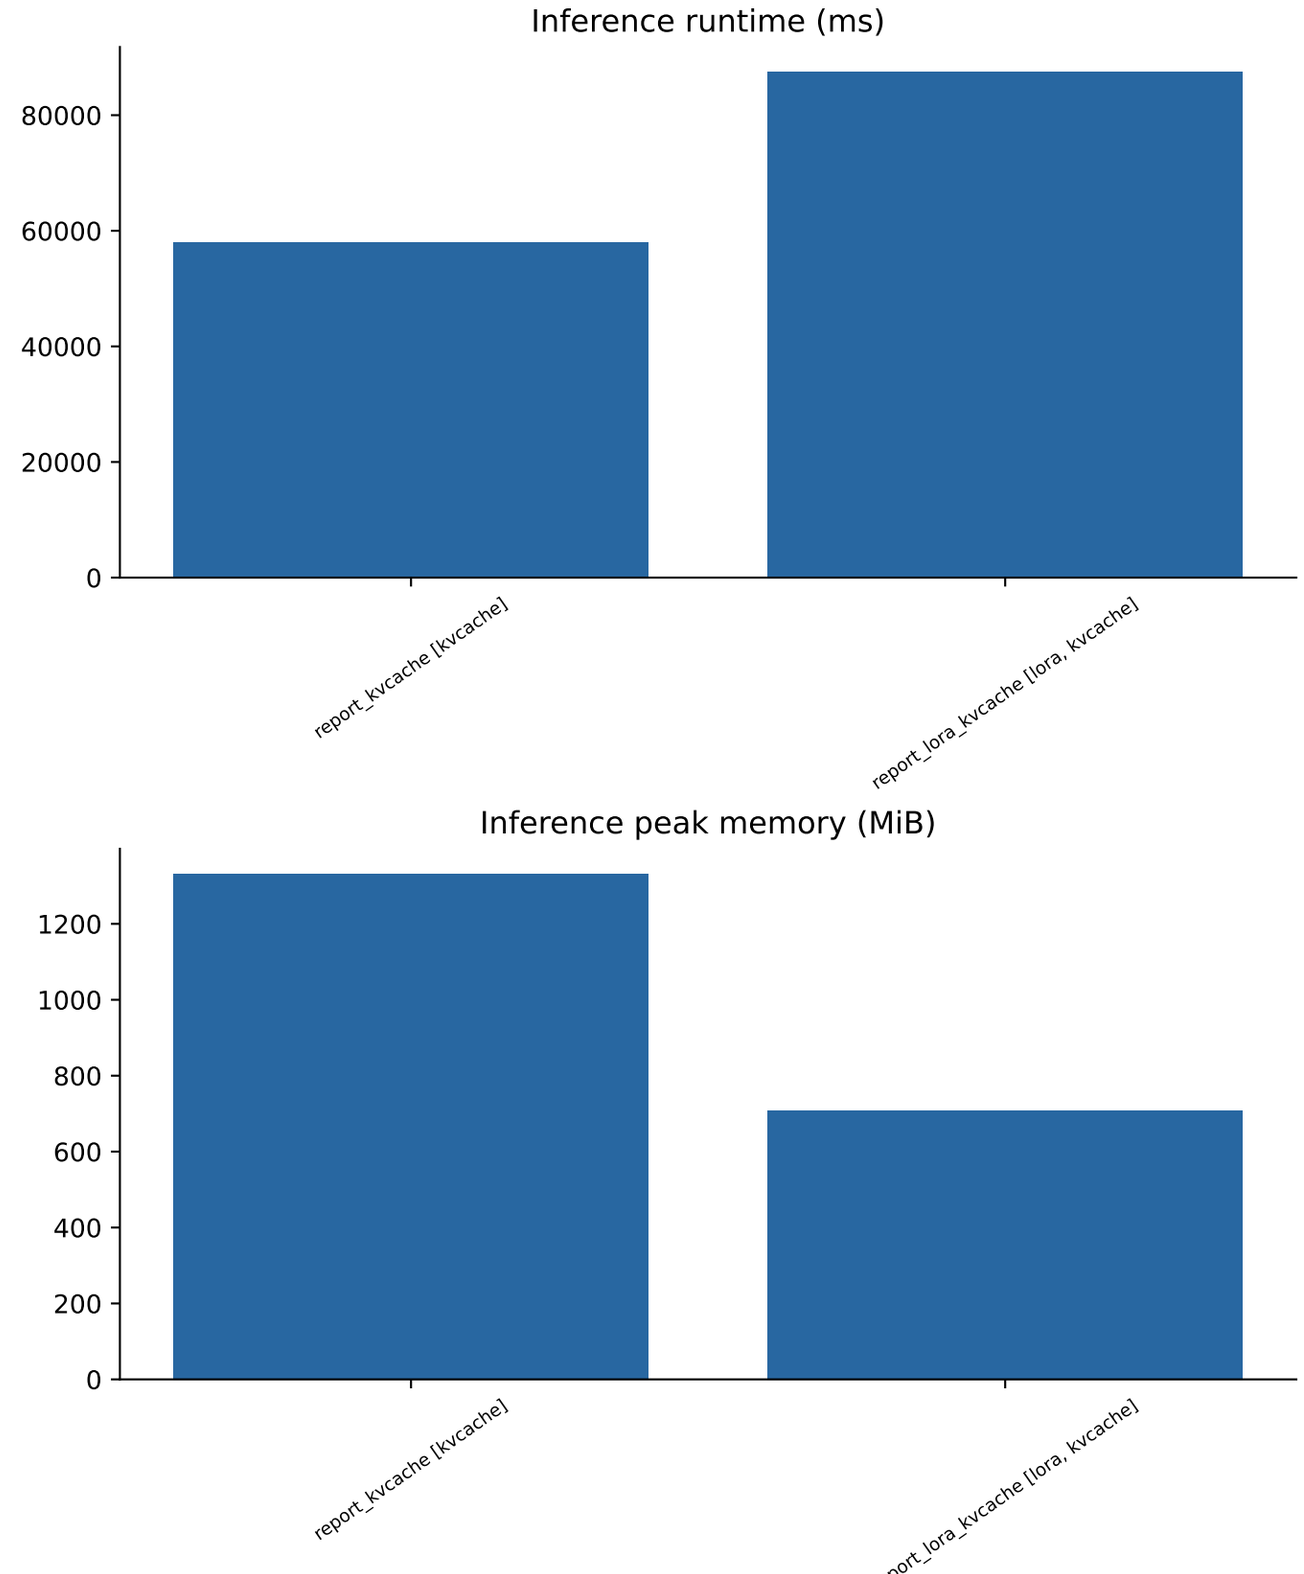
Runtime and memory overview