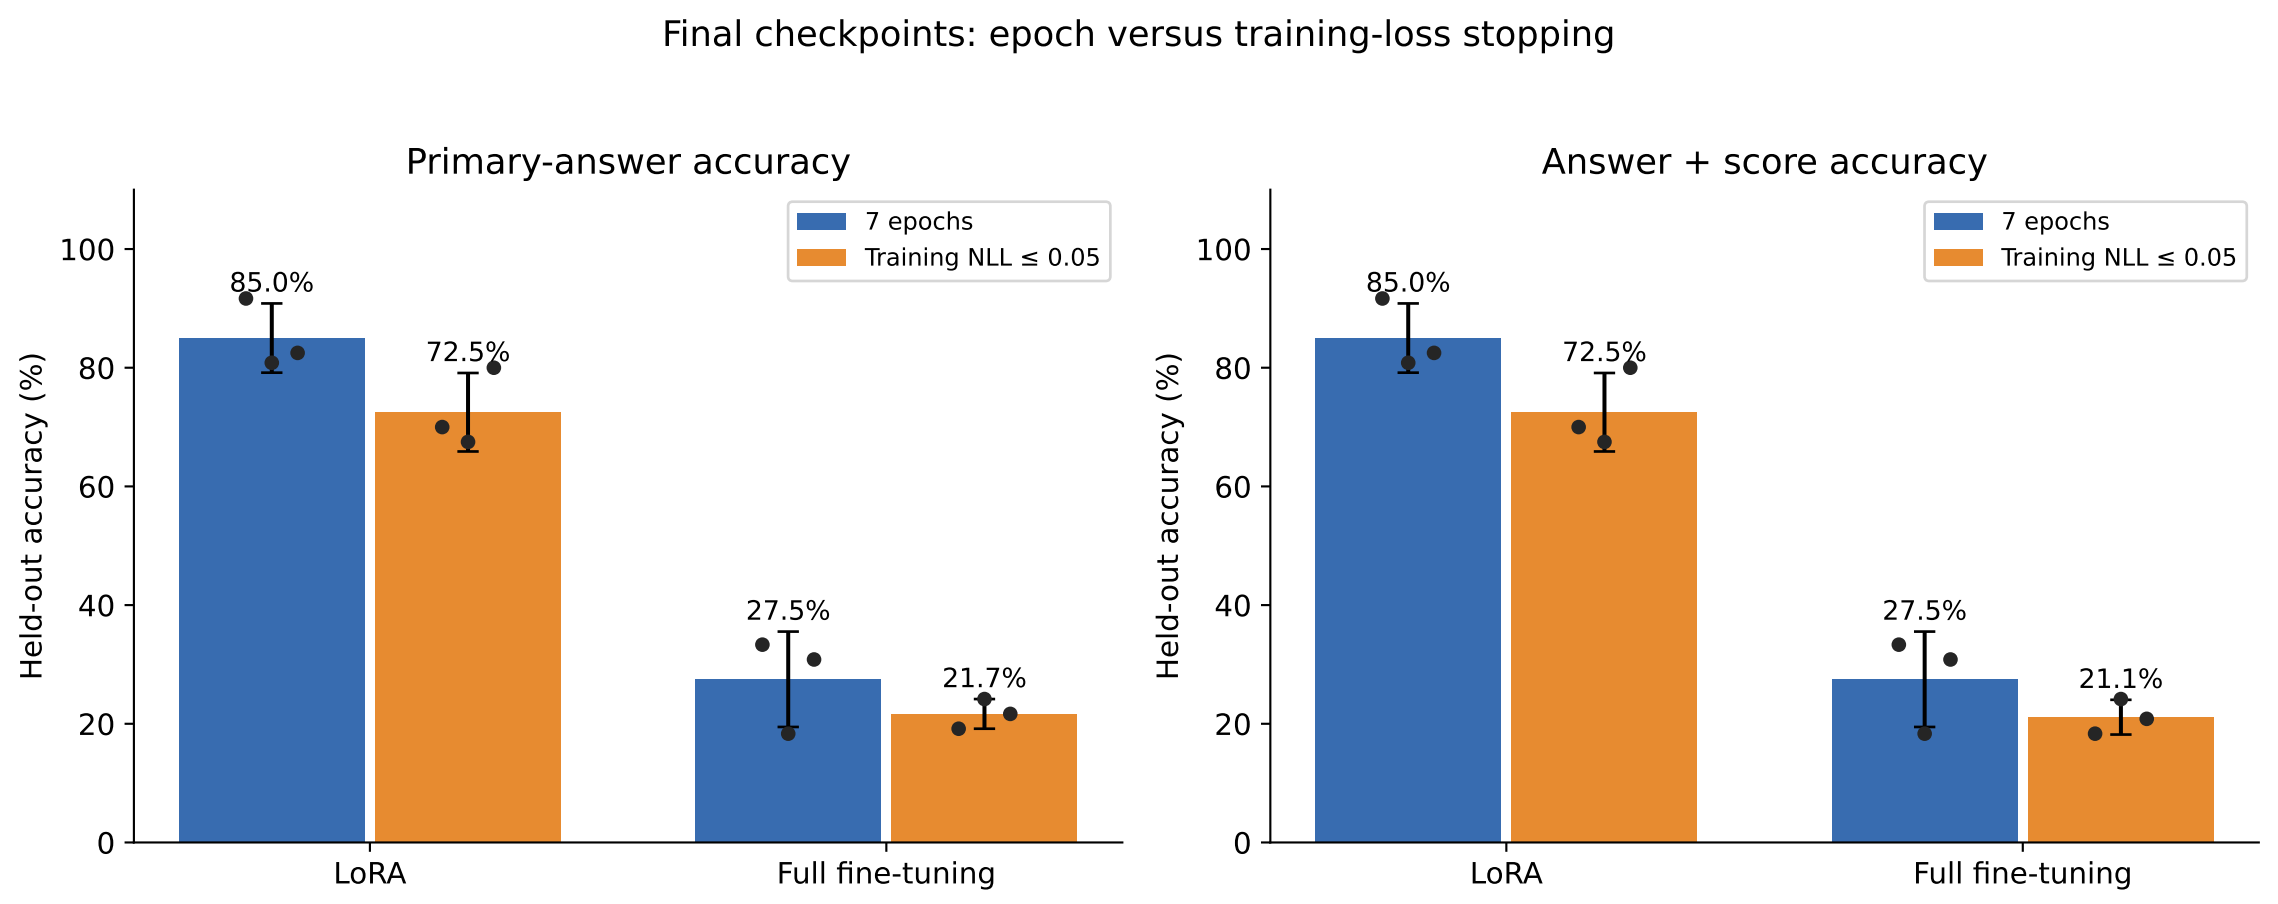
Final checkpoints: epoch versus training-loss stopping
Primary-answer accuracy
Answer + score accuracy
7 epochs
Training NLL ≤ 0.05
Held-out accuracy (%)
0
20
40
60
80
100
85.0%
72.5%
27.5%
21.7%
21.1%
LoRA
Full fine-tuning
Dots: individual seeds. Error bars: sample SD across three seeds. Same existing synthetic benchmark.

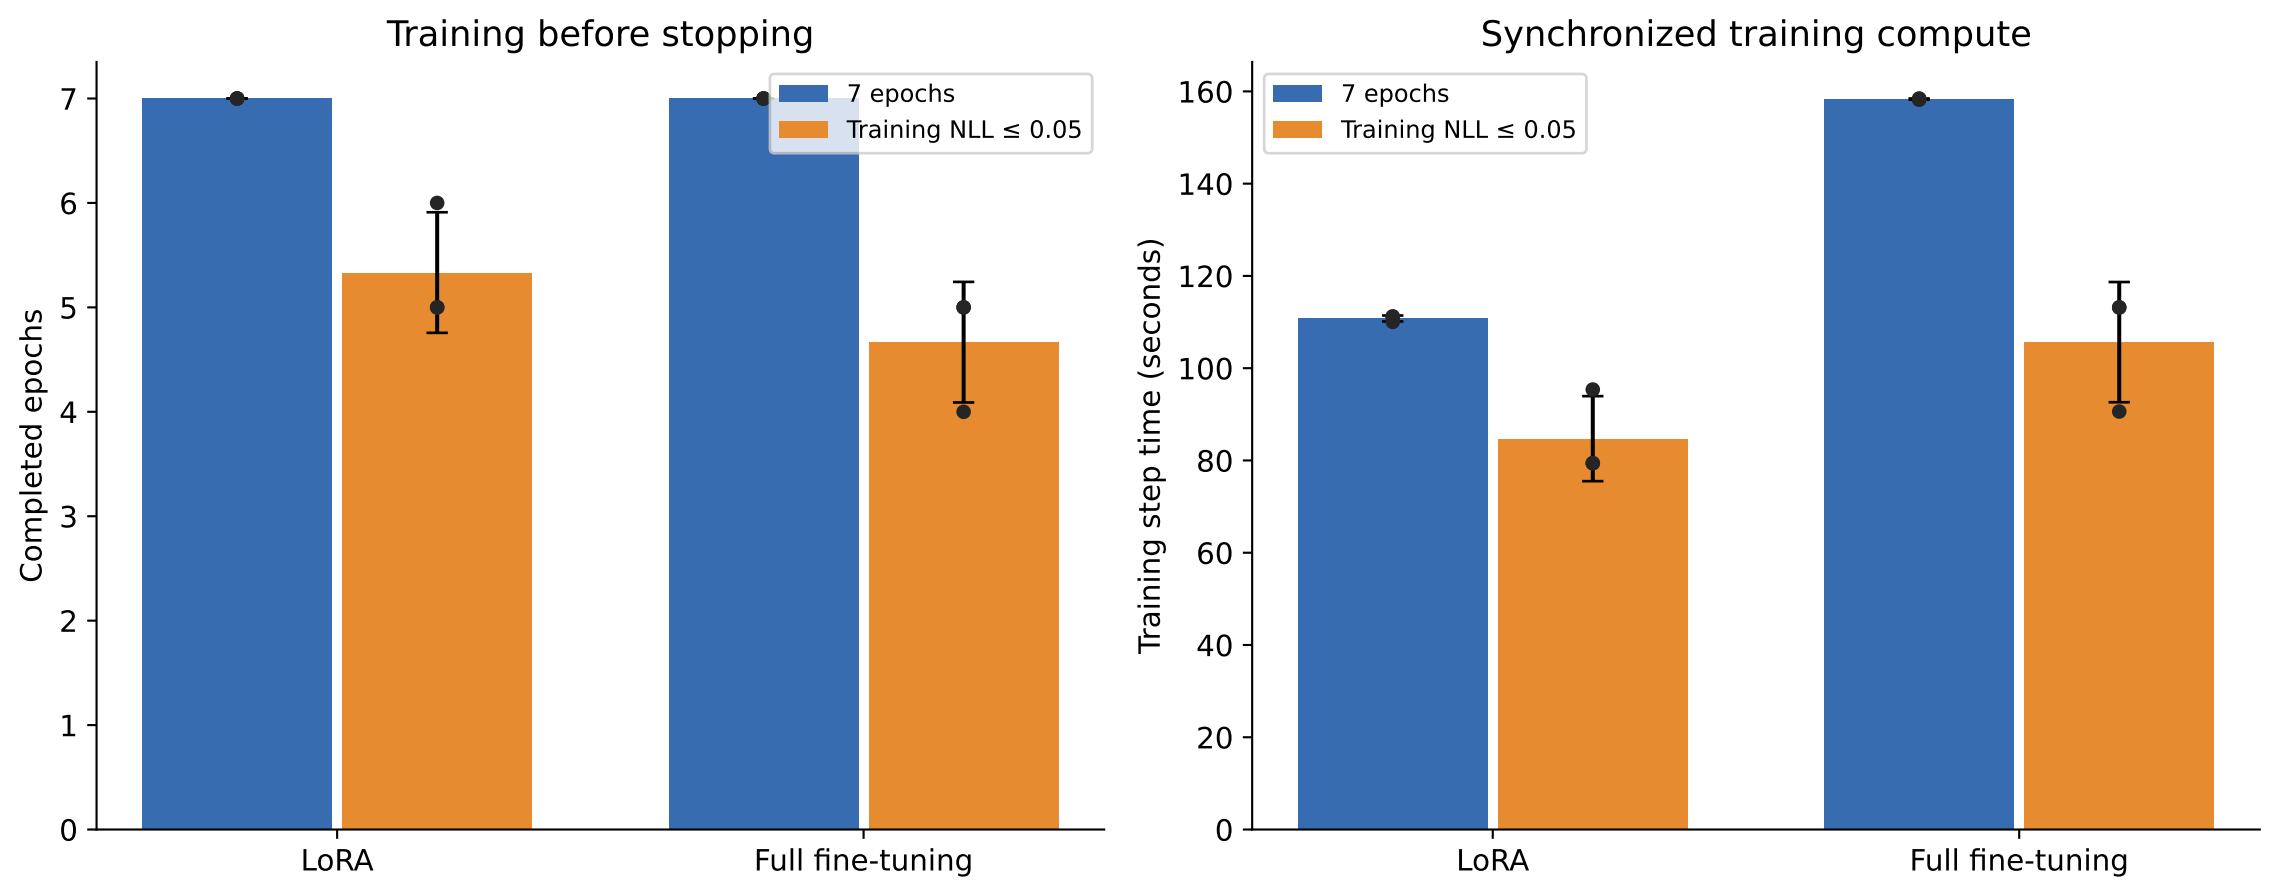

Compute time excludes data transfer, audits, evaluation and checkpoint I/O; it is not end-to-end time.

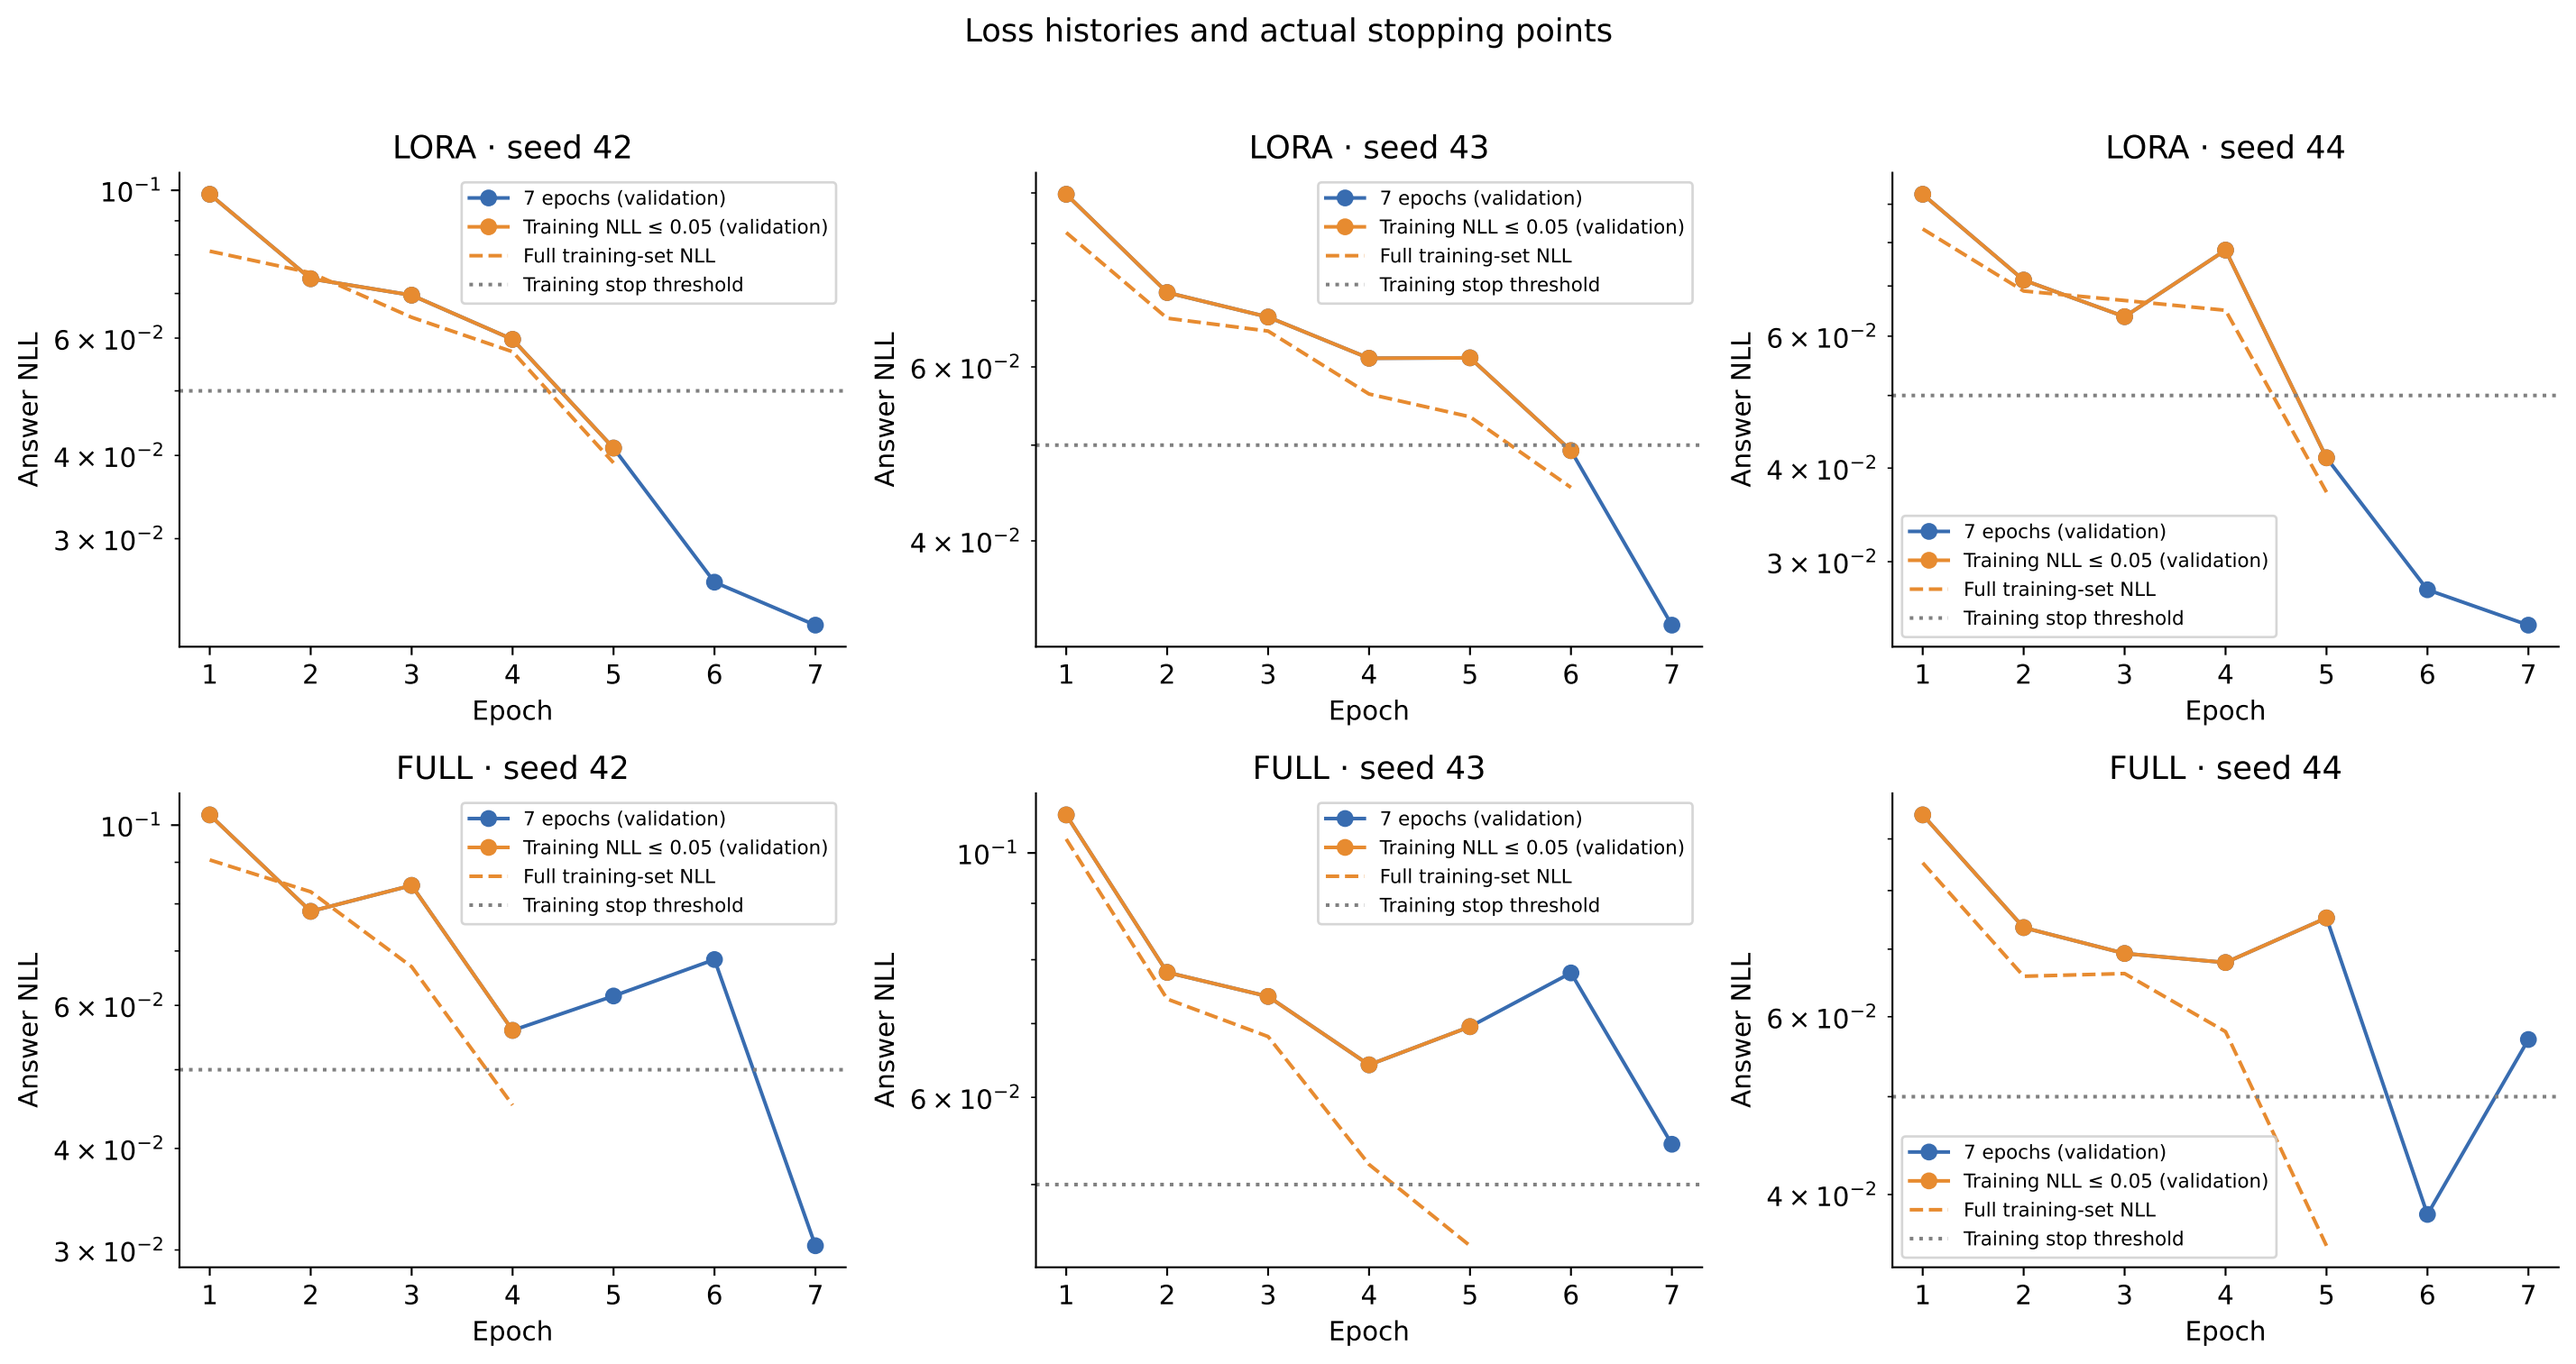
Loss histories and actual stopping points
LORA · seed 42
LORA · seed 43
LORA · seed 44
FULL · seed 42
FULL · seed 43
FULL · seed 44
7 epochs (validation)
Training NLL ≤ 0.05 (validation)
Full training-set NLL
Training stop threshold
Answer NLL
Epoch